

Introduction to ELF

Contents


- What is ELF?
 - Magic number
- ELF Header

What is ELF?

- Executable and Linkable Format
- Common standard file format for
 - Executable file
 - Object code
 - Shared libraries
 - Core dump

What is ELF?

- Made up with of one ELF header, followed by file data
- `readelf` tool can read / display information about ELF files
- The ELF header is 52 bytes long for 32 bit, 64 bytes long for 64 bit addresses



What is ELF?

```
hatchling@DESKTOP-23UPC6C: ~/play
hatchling@DESKTOP-23UPC6C:~/play$ ls
foo1  foo1.c  foo2  foo2.c  output1  output2
hatchling@DESKTOP-23UPC6C:~/play$ readelf -h foo1
ELF Header:
  Magic:   7f 45 4c 46 02 01 01 00 00 00 00 00 00 00 00 00
  Class:                                ELF64
  Data:                                      2's complement, little endian
  Version:                               1 (current)
  OS/ABI:                                UNIX - System V
  ABI Version:                           0
  Type:                                  EXEC (Executable file)
  Machine:                               Advanced Micro Devices X86-64
  Version:                               0x1
  Entry point address:                   0x4003e0
  Start of program headers:              64 (bytes into file)
  Start of section headers:              6384 (bytes into file)
  Flags:                                  0x0
  Size of this header:                   64 (bytes)
  Size of program headers:               56 (bytes)
  Number of program headers:              9
  Size of section headers:               64 (bytes)
  Number of section headers:              28
  Section header string table index:     27
hatchling@DESKTOP-23UPC6C:~/play$
```

ELF header

```
hatchling@DESKTOP-23UPC6C: ~/play
hatchling@DESKTOP-23UPC6C:~/play$ readelf -l foo1

Elf file type is EXEC (Executable file)
Entry point 0x4003e0
There are 9 program headers, starting at offset 64

Program Headers:
  Type           Offset             VirtAddr           PhysAddr
   Type           FileSiz            MemSiz              Flags  Align
  PHDR            0x0000000000000040 0x0000000000400040 0x0000000000400040
  LOAD            0x00000000000001f8 0x00000000000001f8 R      0x8
  INTERP          0x0000000000000238 0x0000000000400238 0x0000000000400238
                  0x00000000000001c 0x00000000000001c  R      0x1
    [Requesting program interpreter: /lib64/ld-linux-x86-64.so.2]
  LOAD            0x00000000000006d0 0x00000000000006d0 R E     0x200000
  LOAD            0x0000000000000e10 0x0000000000000e10 0x0000000000000e10
                  0x0000000000000220 0x0000000000000228 RW     0x200000
  DYNAMIC          0x0000000000000e20 0x0000000000000e20 0x0000000000000e20
                  0x00000000000001d0 0x00000000000001d0 RW     0x8
  NOTE            0x0000000000000254 0x0000000000400254 0x0000000000400254
                  0x0000000000000020 0x0000000000000020 R      0x4
  GNU_EH_FRAME    0x0000000000000594 0x0000000000400594 0x0000000000400594
                  0x000000000000003c 0x000000000000003c R      0x4
  GNU_STACK       0x0000000000000000 0x0000000000000000 0x0000000000000000
                  0x0000000000000000 0x0000000000000000 RW     0x10
  GNU_RELRO       0x0000000000000e10 0x0000000000000e10 0x0000000000000e10
                  0x00000000000001f0 0x00000000000001f0 R      0x1
```

```
hatchling@DESKTOP-23UPC6C: ~/play
hatchling@DESKTOP-23UPC6C:~/play$ readelf -S foo1
There are 28 section headers, starting at offset 0x18f0:

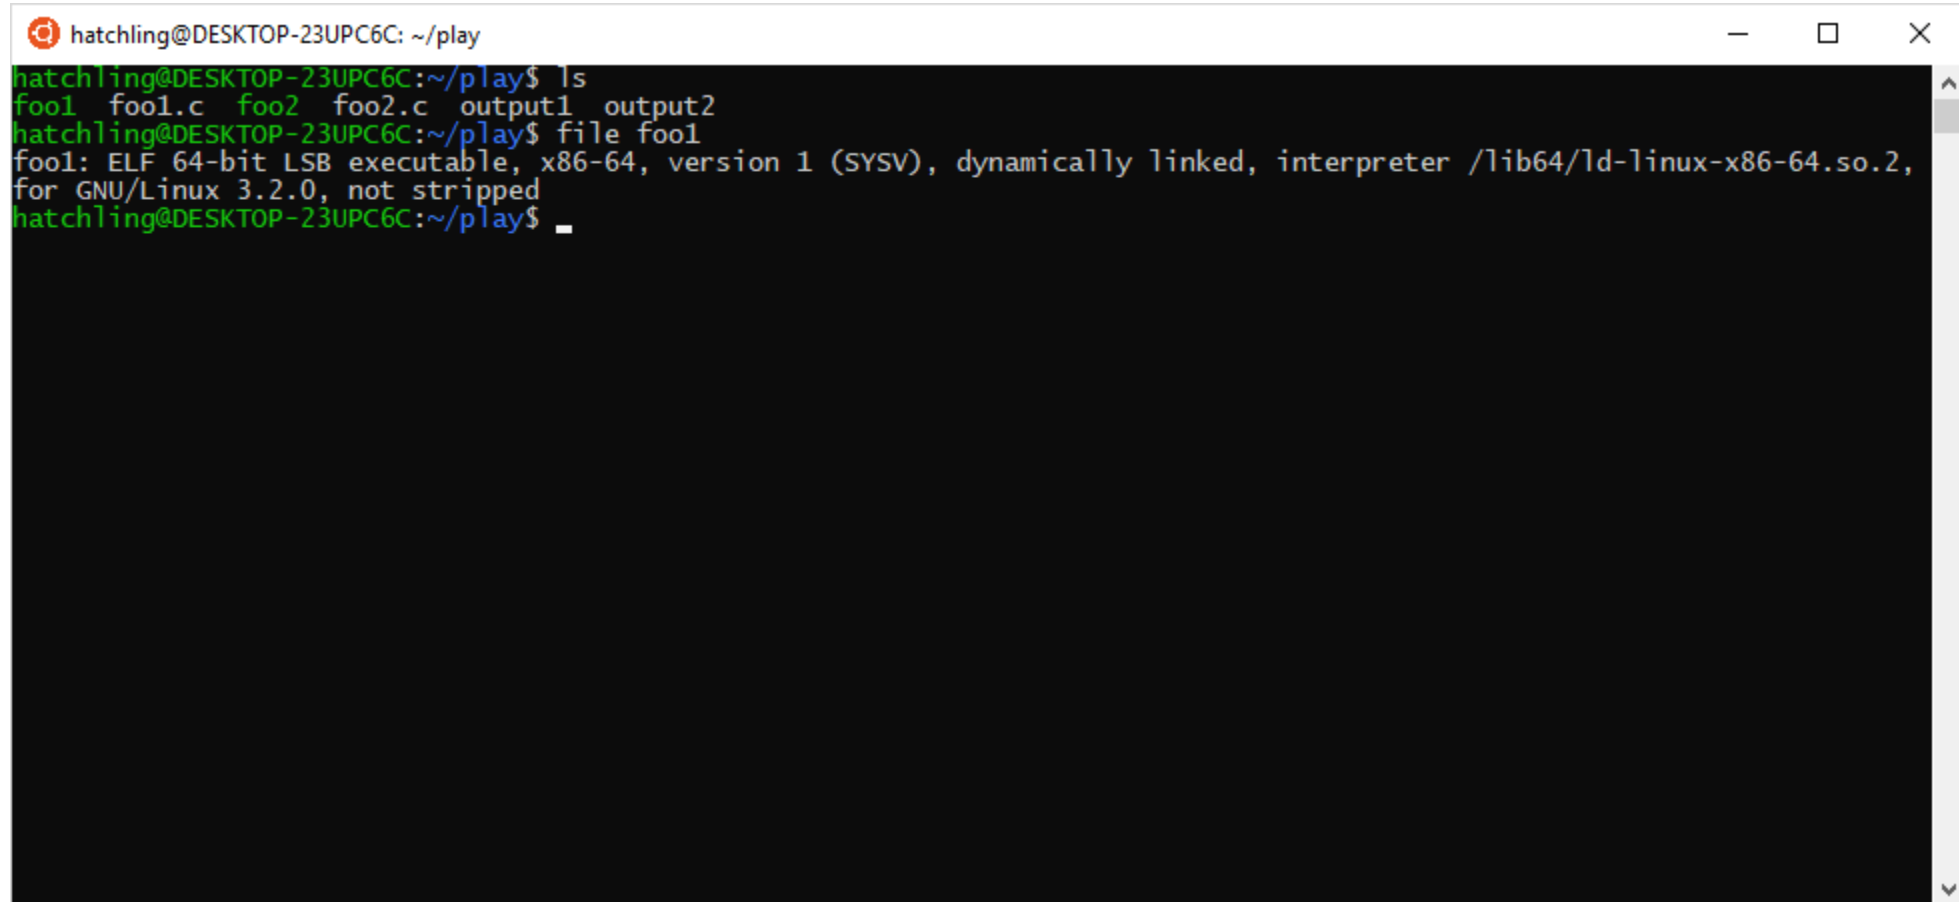
Section Headers:
 [Nr] Name              Type              Address            Offset
     Size            EntSize          Flags  Link  Info  Align
 [ 0]                  NULL              0000000000000000    00000000
     0000000000000000 0000000000000000    0 0
 [ 1] .interp            PROGBITS          0000000000400238    00000238
     000000000000001c 0000000000000000    A 0 0 1
 [ 2] .note.ABI-tag      NOTE              0000000000400254    00000254
     0000000000000020 0000000000000000    A 0 0 4
 [ 3] .gnu.hash          GNU_HASH          0000000000400278    00000278
     000000000000001c 0000000000000000    A 4 0 8
 [ 4] .dynsym            DYNSYM            0000000000400298    00000298
     0000000000000060 0000000000000018    A 5 1 8
 [ 5] .dynstr            STRTAB            00000000004002f8    000002f8
     000000000000003f 0000000000000000    A 0 0 1
 [ 6] .gnu.version        VERSYM            0000000000400338    00000338
     0000000000000008 0000000000000002    A 4 0 2
 [ 7] .gnu.version_r      VERNEED           0000000000400340    00000340
     0000000000000020 0000000000000000    A 5 1 8
 [ 8] .rela.dyn           RELA              0000000000400360    00000360
     0000000000000030 0000000000000018    A 4 0 8
 [ 9] .rela.plt           RELA              0000000000400390    00000390
     0000000000000018 0000000000000018    AI 4 21 8
[10] .init              PROGBITS          00000000004003a8    000003a8
     0000000000000017 0000000000000000    AX 0 0 4
[11] .plt               PROGBITS          00000000004003c0    000003c0
     0000000000000020 0000000000000010    AX 0 0 16
[12] .text              PROGBITS          00000000004003e0    000003e0
     0000000000000192 0000000000000000    AX 0 0 16
```

What is ELF?

Magic number

- Data used to identify the contents of a file
- In Unix-like system, filename extension is not mandatory
- Therefore, we use `file` command to read / interpret magic numbers
- Similar to MIME type

Magic number

A terminal window with a dark background and light-colored text. The window title bar shows 'hatchling@DESKTOP-23UPC6C: ~/play' and standard window controls. The terminal shows the execution of 'ls' and 'file' commands. The 'file' command output identifies 'foo1' as an ELF 64-bit LSB executable for x86-64 architecture, dynamically linked with the GNU/Linux 3.2.0 interpreter.

```
hatchling@DESKTOP-23UPC6C: ~/play
hatchling@DESKTOP-23UPC6C:~/play$ ls
foo1  foo1.c  foo2  foo2.c  output1  output2
hatchling@DESKTOP-23UPC6C:~/play$ file foo1
foo1: ELF 64-bit LSB executable, x86-64, version 1 (SYSV), dynamically linked, interpreter /lib64/ld-linux-x86-64.so.2,
for GNU/Linux 3.2.0, not stripped
hatchling@DESKTOP-23UPC6C:~/play$
```


What is ELF?

- Program header table
 - Tells the system how to create a process image:
 - A copy of the process at a given point in time
 - Files used to execute a program must have one
- Section header table
 - Contains information describing the file's sections
 - Files used during linking must have one

ELF header

```
hatchling@DESKTOP-23UPC6C: ~/play
hatchling@DESKTOP-23UPC6C:~/play$ ls
foo1  foo1.c  foo2  foo2.c  output1  output2
hatchling@DESKTOP-23UPC6C:~/play$ readelf -h foo1
ELF Header:
  Magic:   7f 45 4c 46 02 01 01 00 00 00 00 00 00 00 00 00
  Class:                             ELF64
  Data:                               2's complement, little endian
  Version:                           1 (current)
  OS/ABI:                             UNIX - System V
  ABI Version:                       0
  Type:                               EXEC (Executable file)
  Machine:                           Advanced Micro Devices X86-64
  Version:                           0x1
  Entry point address:                0x4003e0
  Start of program headers:          64 (bytes into file)
  Start of section headers:         6384 (bytes into file)
  Flags:                              0x0
  Size of this header:                64 (bytes)
  Size of program headers:           56 (bytes)
  Number of program headers:          9
  Size of section headers:           64 (bytes)
  Number of section headers:         28
  Section header string table index: 27
hatchling@DESKTOP-23UPC6C:~/play$
```

```
hatchling@DESKTOP-23UPC6C: ~/play
hatchling@DESKTOP-23UPC6C:~/play$ readelf -h fool
ELF Header:
  Magic:   7f 45 4c 46 02 01 01 00 00 00 00 00 00 00 00 00
  Class:                                ELF64
  Data:                                      2's complement, little endian
  Version:                               1 (current)
  OS/ABI:                                UNIX - System V
  ABI Version:                           0
  Type:                                  EXEC (Executable file)
  Machine:                               Advanced Micro Devices X86-64
  Version:                               0x1
  Entry point address:                   0x4003e0
  Start of program headers:              64 (bytes into file)
  Start of section headers:              6384 (bytes into file)
  Flags:                                  0x0
  Size of this header:                   64 (bytes)
  Size of program headers:               56 (bytes)
  Number of program headers:              9
  Size of section headers:               64 (bytes)
  Number of section headers:             28
  Section header string table index:     27
hatchling@DESKTOP-23UPC6C:~/play$
```

```
hatchling@DESKTOP-23UPC6C: ~/play
00000000: 7f45 4c46 0201 0100 0000 0000 0000 0000 .ELF.....
00000010: 0200 3e00 0100 0000 e003 4000 0000 0000 ..>.....@....
00000020: 4000 0000 0000 0000 f018 0000 0000 0000 @.....8....@...
00000030: 0000 0000 4000 3800 0900 4000 1c00 1b00 .....@.8....@...
00000040: 0600 0000 0400 0000 4000 0000 0000 0000 .....@.....
00000050: 4000 4000 0000 0000 4000 4000 0000 0000 @.@.....@.@....
00000060: f801 0000 0000 0000 f801 0000 0000 0000 .....
00000070: 0800 0000 0000 0000 0300 0000 0400 0000 8.....8.@....
00000080: 3802 0000 0000 0000 3802 4000 0000 0000 8.@.....
00000090: 3802 4000 0000 0000 1c00 0000 0000 0000 8.@.....
000000a0: 1c00 0000 0000 0000 0100 0000 0000 0000 .....
000000b0: 0100 0000 0500 0000 0000 0000 0000 0000 .....
000000c0: 0000 4000 0000 0000 0000 4000 0000 0000 ..@.....@....
000000d0: d006 0000 0000 0000 d006 0000 0000 0000 .....
000000e0: 0000 2000 0000 0000 0100 0000 0600 0000 .....
000000f0: 100e 0000 0000 0000 100e 6000 0000 0000 .....
00000100: 100e 6000 0000 0000 2002 0000 0000 0000 .....
00000110: 2802 0000 0000 0000 0000 2000 0000 0000 (.....
00000120: 0200 0000 0600 0000 200e 0000 0000 0000 .....
00000130: 200e 6000 0000 0000 200e 6000 0000 0000 .....
00000140: d001 0000 0000 0000 d001 0000 0000 0000 .....
00000150: 0800 0000 0000 0000 0400 0000 0400 0000 .....
00000160: 5402 0000 0000 0000 5402 4000 0000 0000 T.....T.@....
00000170: 5402 4000 0000 0000 2000 0000 0000 0000 T.@.....
00000180: 2000 0000 0000 0000 0400 0000 0000 0000 .....
00000190: 50e5 7464 0400 0000 9405 0000 0000 0000 P.td.....
000001a0: 9405 4000 0000 0000 9405 4000 0000 0000 ..@.....@....
000001b0: 3c00 0000 0000 0000 3c00 0000 0000 0000 <.....<.....
000001c0: 0400 0000 0000 0000 51e5 7464 0600 0000 .....Q.td....
fool.x
```

ELF header

Overall

- ELF is common standard file format used in Unix-like system
- Made up of one ELF header, followed by file data
- Magic number is [0x7F, 'E'(0x45), 'L'(0x4C), 'F'(0x46)]
- Difference between section vs. segment?

Citation

- <https://refspecs.linuxfoundation.org/elf/elf.pdf>
- (1) - <https://en.wikipedia.org/wiki/File:Elf-layout--en.svg>